

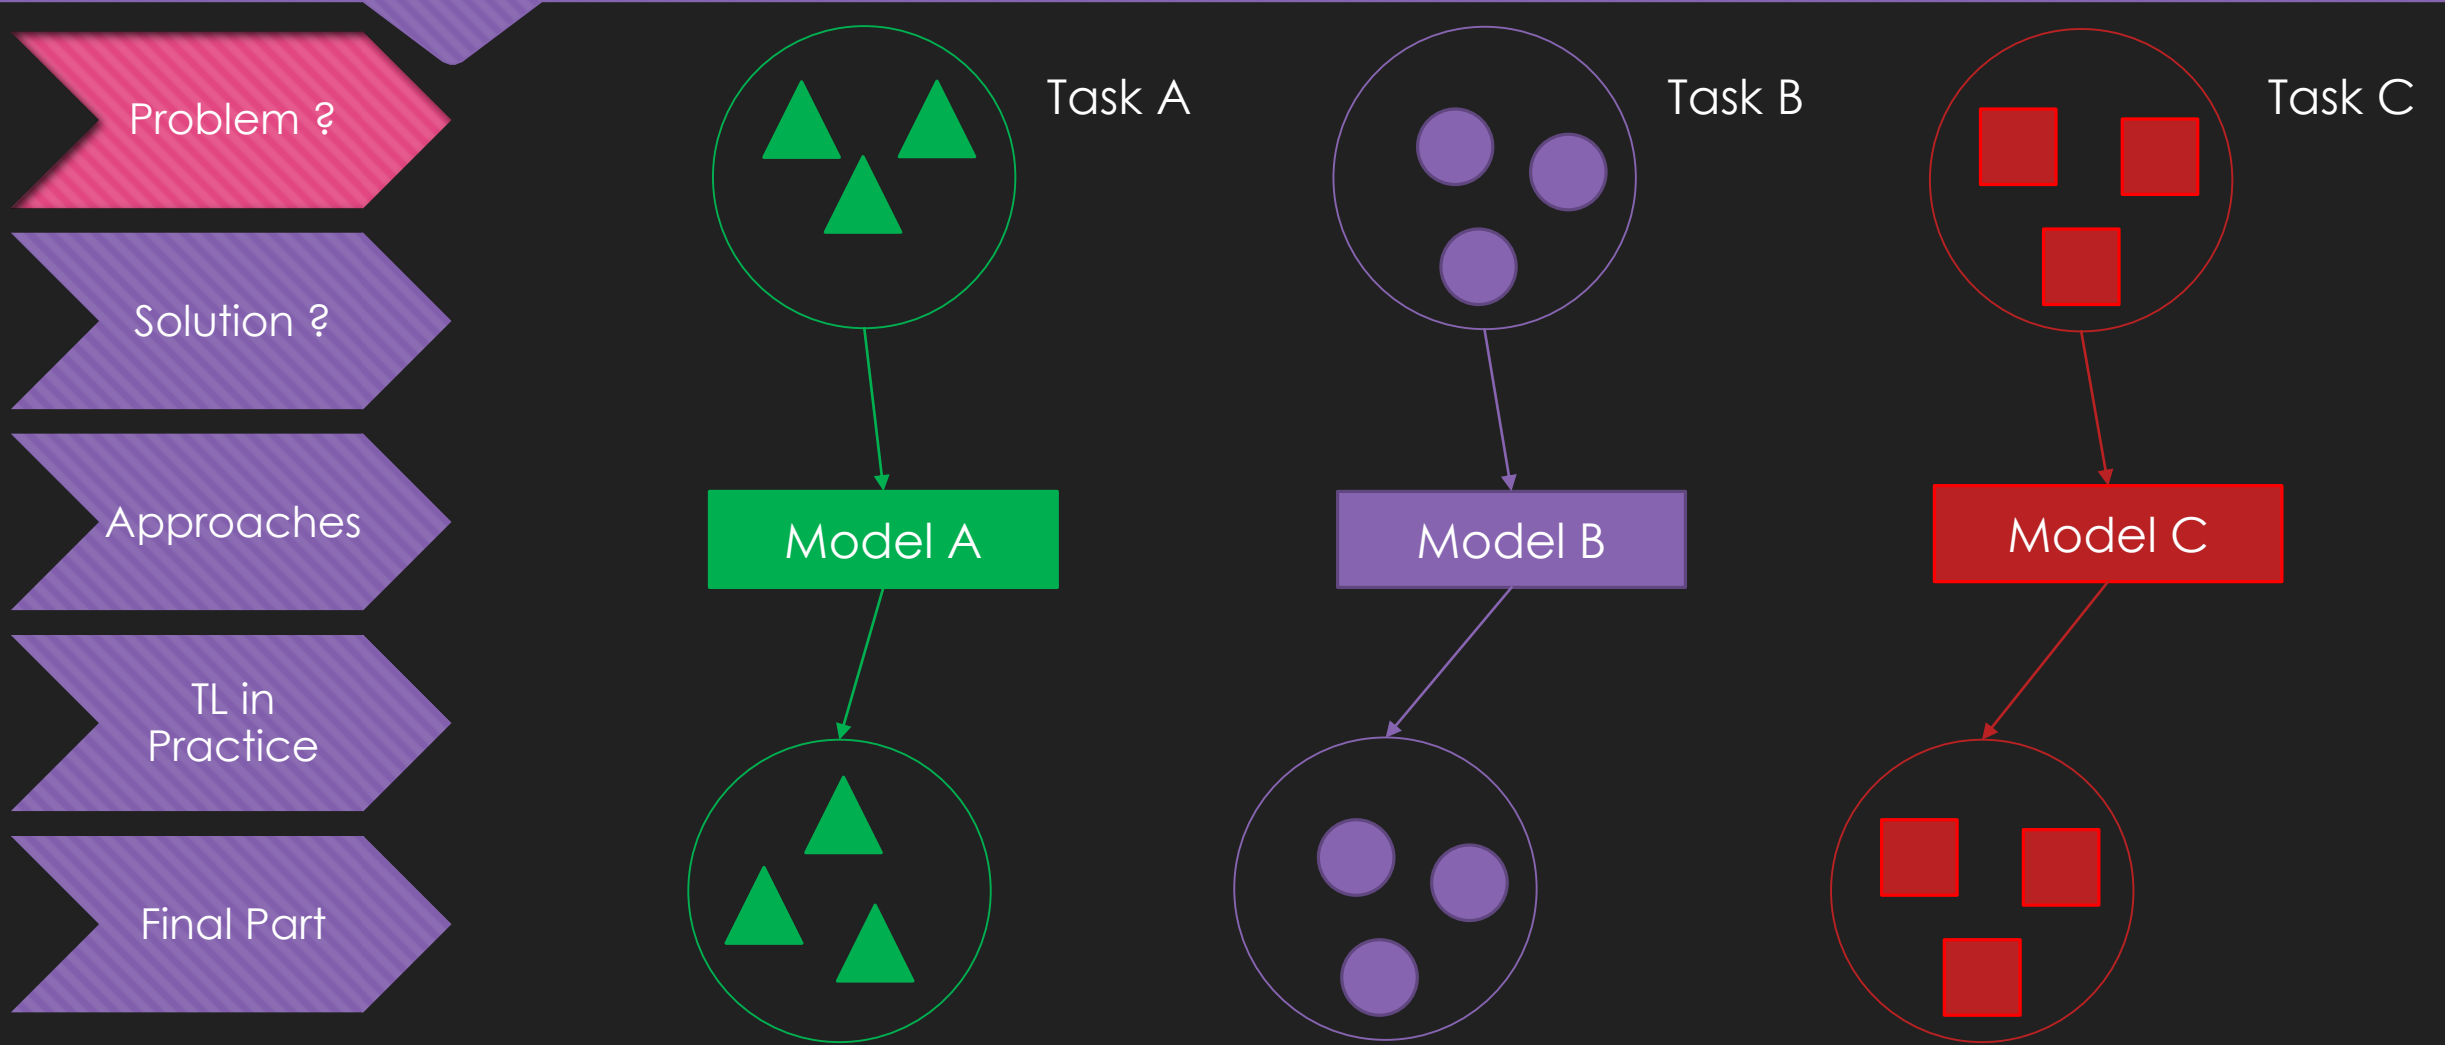
Transfer Learning Theory and it's Applications using Tensorflow & Keras

Babak Badnava

PyCon 2018



Traditional ML



Problem Definition!

Problem ?

Solution ?

Approaches

TL in
Practice

Final Part

- Complex problems
- Insufficient labeled data
- Long training time

Solution ?

Problem ?

- Assumptions ?

Solution ?

- Related data for different task

- Previously learned tasks

Approaches

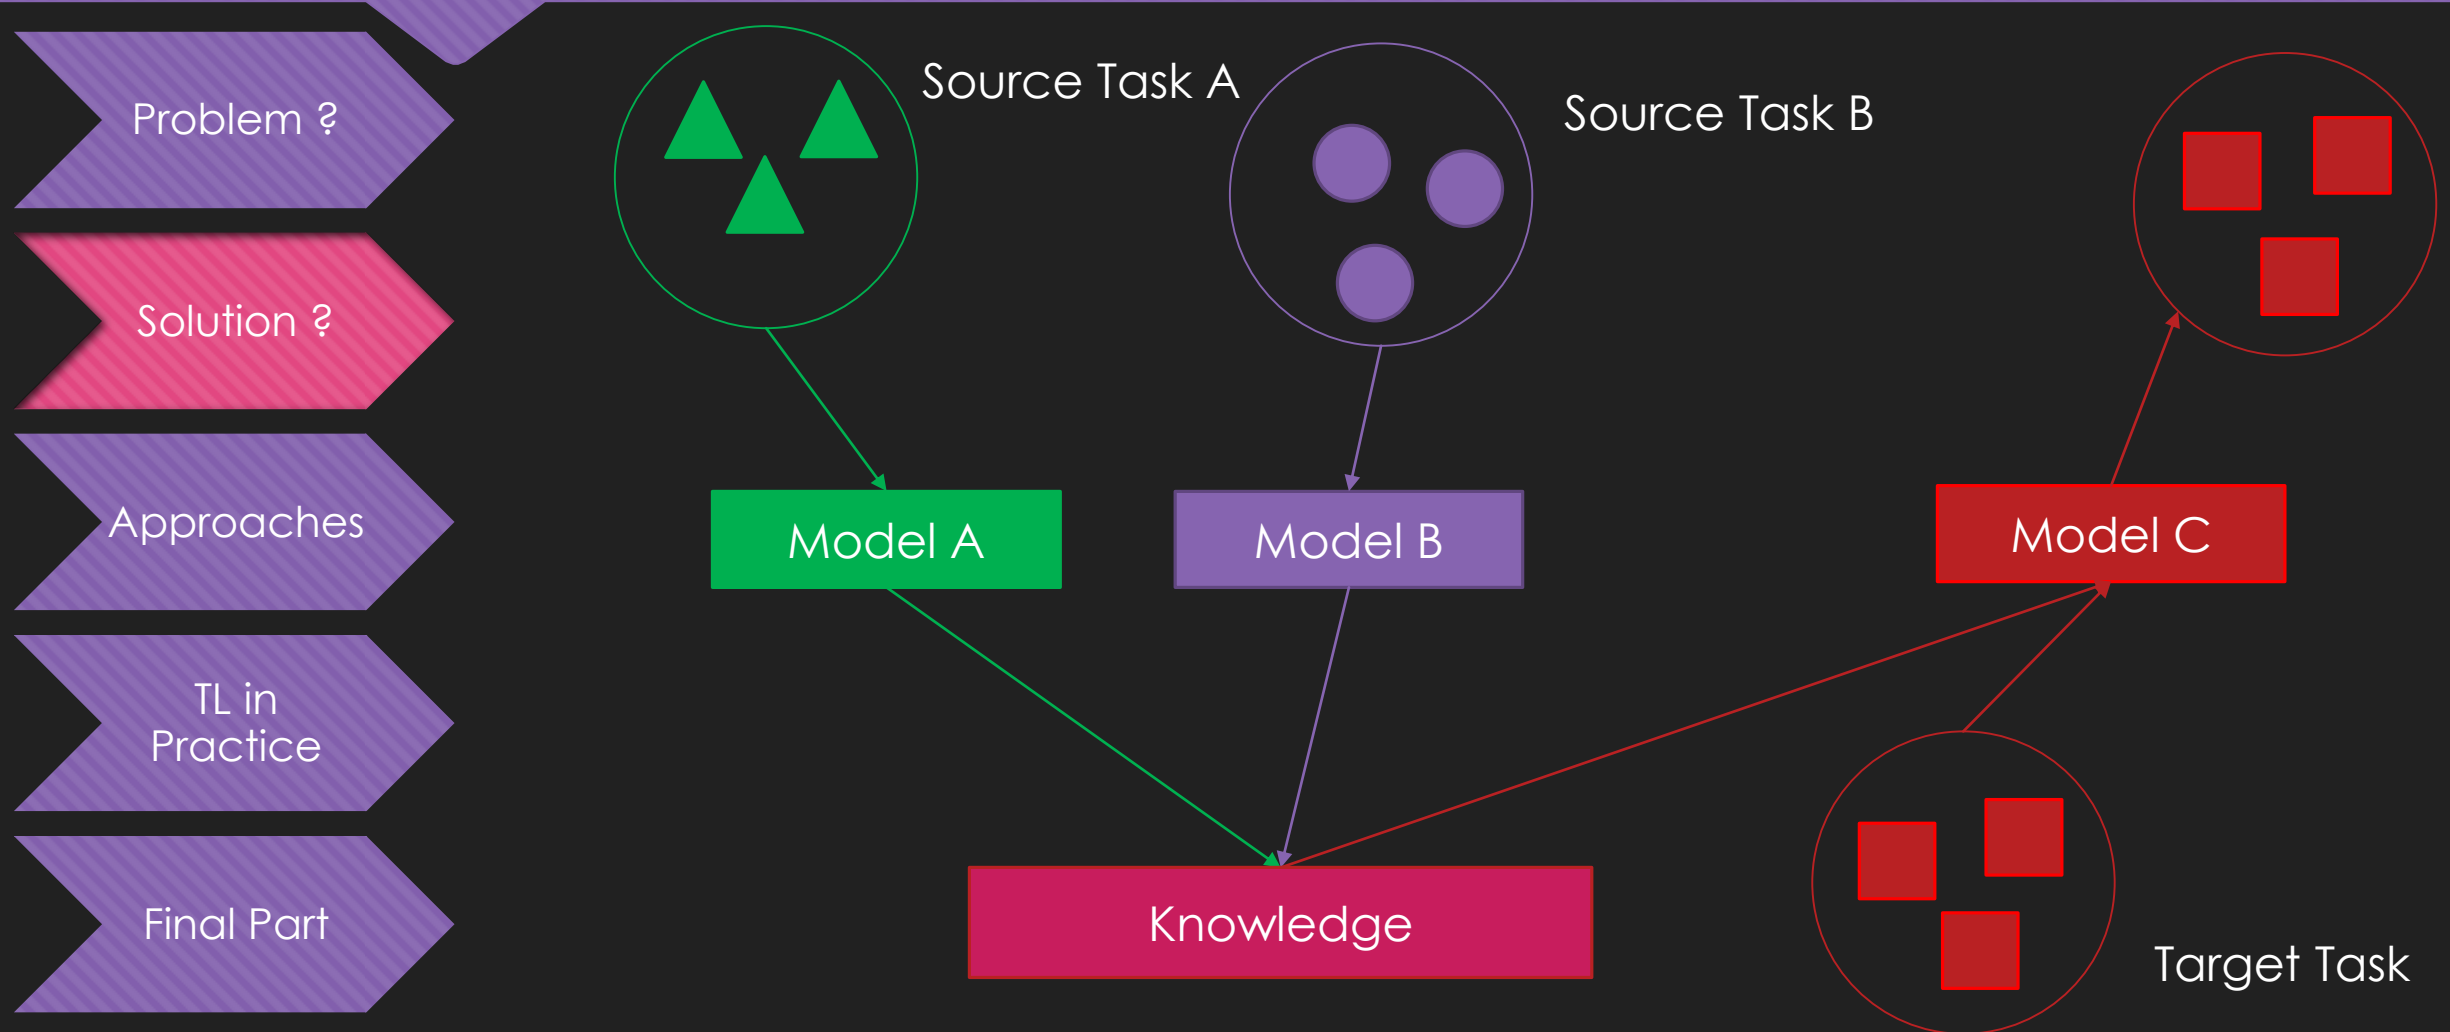
- Solution ?

TL in
Practice

- Use them

Final Part

Transfer Learning framework



Transfer Learning and Neural Networks

Problem ?

- Pre-trained models

Solution ?

- Trains with large dataset
- Use for specific task

Approaches

- Fixed feature extractor

TL in
Practice

- Autoencoders
- Freezing pre-trained model

Final Part

- Fine-tuning

Pre-trained models on ImageNet in Keras

Problem ?

○ Xception

Solution ?

○ VGG16

○ VGG19

Approaches

○ ResNet50

○ InceptionV3

TL in
Practice

○

Final Part

<https://keras.io/applications/>

ResNet50 for specific task

Problem ?

Solution ?

Approaches

TL in
Practice

Final Part

```
from keras.applications.resnet50 import ResNet50
from keras.preprocessing import image
from keras.applications.resnet50 import preprocess_input, decode_predictions
import numpy as np

model = ResNet50(weights='imagenet')

img_path = 'elephant.jpg'
img = image.load_img(img_path, target_size=(224, 224))
x = image.img_to_array(img)
x = np.expand_dims(x, axis=0)
x = preprocess_input(x)

preds = model.predict(x)
```


Extract features with VGG16

Problem ?

Solution ?

Approaches

TL in
Practice

Final Part

```
from keras.applications.vgg16 import VGG16
from keras.preprocessing import image
from keras.applications.vgg16 import preprocess_input
import numpy as np

model = VGG16(weights='imagenet', include_top=False)

img_path = 'elephant.jpg'
img = image.load_img(img_path, target_size=(224, 224))
x = image.img_to_array(img)
x = np.expand_dims(x, axis=0)
x = preprocess_input(x)

features = model.predict(x)
```

Extract features from an intermediate layer with VGG19

Problem ?

Solution ?

Approaches

TL in
Practice

Final Part

```
from keras.applications.vgg19 import VGG19
from keras.preprocessing import image
from keras.applications.vgg19 import preprocess_input
from keras.models import Model
import numpy as np

base_model = VGG19(weights='imagenet')
model = Model(inputs=base_model.input, outputs=base_model.get_layer('block4_pool').output)

img_path = 'elephant.jpg'
img = image.load_img(img_path, target_size=(224, 224))
x = image.img_to_array(img)
x = np.expand_dims(x, axis=0)
x = preprocess_input(x)

block4_pool_features = model.predict(x)
```

Fine-tune InceptionV3 on a new set of classes

Problem ?

Solution ?

Approaches

TL in
Practice

Final Part

```
from keras.applications.inception_v3 import InceptionV3
from keras.preprocessing import image
from keras.models import Model
from keras.layers import Dense, GlobalAveragePooling2D
from keras import backend as K

# create the base pre-trained model
base_model = InceptionV3(weights='imagenet', include_top=False)

# add a global spatial average pooling layer
x = base_model.output
x = GlobalAveragePooling2D()(x)
# let's add a fully-connected layer
x = Dense(1024, activation='relu')(x)
# and a logistic layer -- let's say we have 200 classes
predictions = Dense(200, activation='softmax')(x)

model = Model(inputs=base_model.input, outputs=predictions)
```

Fine-tune InceptionV3 on a new set of classes(2)

Problem ?

Solution ?

Approaches

TL in
Practice

Final Part

```
for layer in base_model.layers:
    layer.trainable = False

model.compile(optimizer='rmsprop', loss='categorical_crossentropy')

model.fit_generator(...)

for layer in model.layers[:249]:
    layer.trainable = False
for layer in model.layers[249:]:
    layer.trainable = True

from keras.optimizers import SGD
model.compile(optimizer=SGD(lr=0.0001, momentum=0.9), loss='categorical_crossentropy')

model.fit_generator(...)
```

TensorFlow-Slim

Problem ?

Solution ?

Approaches

TL in
Practice

Final Part

- lightweight high-level API of TensorFlow

- Many pre-trained models such as:

- Inception

- ResNet V1 50

- ResNet V1 101

- VGG 16

<https://github.com/tensorflow/models/tree/master/research/slim>

Pre-trained models

Problem ?

Solution ?

Approaches

TL in
Practice

Final Part

```
import tensorflow as tf
import tensorflow.contrib.slim.nets as nets

slim = tf.contrib.slim
vgg = nets.vgg
train_log_dir = ...
if not tf.gfile.Exists(train_log_dir):
    tf.gfile.MakeDirs(train_log_dir)

with tf.Graph().as_default():
    # Set up the data loading:
    images, labels = ...
    # Define the model:
    predictions = vgg.vgg_16(images, is_training=True)
    # Specify the loss function:
    slim.losses.softmax_cross_entropy(predictions, labels)
    total_loss = slim.losses.get_total_loss()
    tf.summary.scalar('losses/total_loss', total_loss)
    # Specify the optimization scheme:
    optimizer = tf.train.GradientDescentOptimizer(learning_rate=.001)
    train_tensor = slim.learning.create_train_op(total_loss, optimizer)

    # Actually runs training.
    slim.learning.train(train_tensor, train_log_dir)
```


Transfer a specific layer

Problem ?

Solution ?

Approaches

TL in
Practice

Final Part

```
import tensorflow as tf
```

```
exclude = ['InceptionResnetV2/Logits', 'InceptionResnetV2/AuxLogits']  
variables_to_restore = tf.contrib.slim.get_variables_to_restore(exclude = exclude)  
saver = tf.train.Saver(variables_to_restore)  
checkpoint_file = './inception_resnet_v2_2016_08_30.ckpt'  
saver.restore(sess, checkpoint_file)
```

Freeze a layer

Problem ?

Solution ?

Approaches

TL in
Practice

Final Part

```
def dense_layer(self, input, out_dim, name, func=tf.nn.relu):
    in_dim = input.get_shape().as_list()[-1]
    d = 1.0 / np.sqrt(in_dim)
    with tf.variable_scope(name):
        w_init = tf.random_uniform_initializer(-d, d)
        b_init = tf.random_uniform_initializer(-d, d)
        w = tf.get_variable('w', dtype=tf.float32, shape=[in_dim, out_dim], initializer=w_init, trainable=False)
        b = tf.get_variable('b', shape=[out_dim], initializer=b_init, trainable=False)

        output = tf.matmul(input, w) + b
        if func is not None:
            output = func(output)

    return output
```


Freeze a layer

Problem ?

Solution ?

Approaches

TL in
Practice

Final Part

```
def conv2d_layer(self, input, filter_size, out_dim, name, strides, func=tf.nn.relu):
    in_dim = input.get_shape().as_list()[-1]
    d = 1.0 / np.sqrt(filter_size * filter_size * in_dim)
    with tf.variable_scope(name):
        w_init = tf.random_uniform_initializer(-d, d)
        b_init = tf.random_uniform_initializer(-d, d)
        w = tf.get_variable('w',
                            shape=[filter_size, filter_size, in_dim, out_dim],
                            dtype=tf.float32,
                            initializer=w_init, trainable=False)
        b = tf.get_variable('b', shape=[out_dim], initializer=b_init, trainable=False)

        output = tf.nn.conv2d(input, w, strides=strides, padding='SAME') + b
        if func is not None:
            output = func(output)

    return output
```

Applications

Problem ?

- Learning from simulations

Solution ?

- Autonomous cars

- Robotics

Approaches

- Adapting to new domains

TL in
Practice

- Semi supervised learning

Final Part

When and which?

Problem ?

Solution ?

Approaches

TL in
Practice

Final Part

	Similar dataset	Different dataset
Small dataset	Highest level features + classifier	Lower level features + classifier
Large dataset	Fine tune	Fine tune

Thanks for attention.

Questions?